

Operating Systems: File Management

File management

- Why talk about file management?
 - Files are the central element to most applications
 - The File System is one of the most important parts of the OS to the user
 - Some users tend to think the OS is ONLY the file system!

File management

- Desirable properties of file (systems):
 - Long-term existence
 - Must be able to store very large amounts of information
 - Sharable between processes
 - Want way to information in a way that is easy to access
 - We will need structure
 - Guarantee that data in the file is valid
 - Minimize lost or destroyed data

File management

- Basic functions of a file system
 - Present a logical (abstract) view of files to users by hiding physical details
 - Facilitate the sharing of physical I/O devices
 - Optimize the usage of these I/O devices
 - Provide protection mechanisms for data being transferred or managed by I/O devices

File management

- File management system consists of system utility programs that run as privileged applications
- Concerned with secondary storage
- Usually provide functions like:
 - Create
 - Delete
 - Open
 - Close
 - Read
 - Write

File management

- Provides services to users and applications in the use of files
- File structure at user level usually given by the following terms:
 - Field
 - Record
 - File
 - Database

Fields and Records

- Fields
 - Basic element of data
 - Contains a single value (such as name or date)
 - May be fixed length or variable
 - If variable, separated by demarcation fields
- Records
 - Collection of related fields, can be fixed or variable length
 - Has a field with a length attribute
 - Treated as a unit
 - Examples: Employee record

Files

- File
 - Collection of similar records
 - Treated as a single entity by users and applications
 - So has a filename
 - Usually the smallest unit with access control restrictions
 - Though this can be down to individual fields

Database

- Collection of related data
- Explicit relationships among data elements and fields
- May contain one or more files
- May be managed by a database management system
 - In other words, could be independent of OS

Requirements for general purpose system

- Each user should be able to create, delete, read, write and modify files
- Each user may have controlled access to other users' files
- Each user may control what type of accesses are allowed to the users' files
- Each user should be able to set up their own structure of files
 - In a form appropriate to their needs
- Each user should be able to move data between files
- Each user should be able to back up and restore data in case of damage
- Should provide a convenient method to access files (symbolic names)

Typical organization

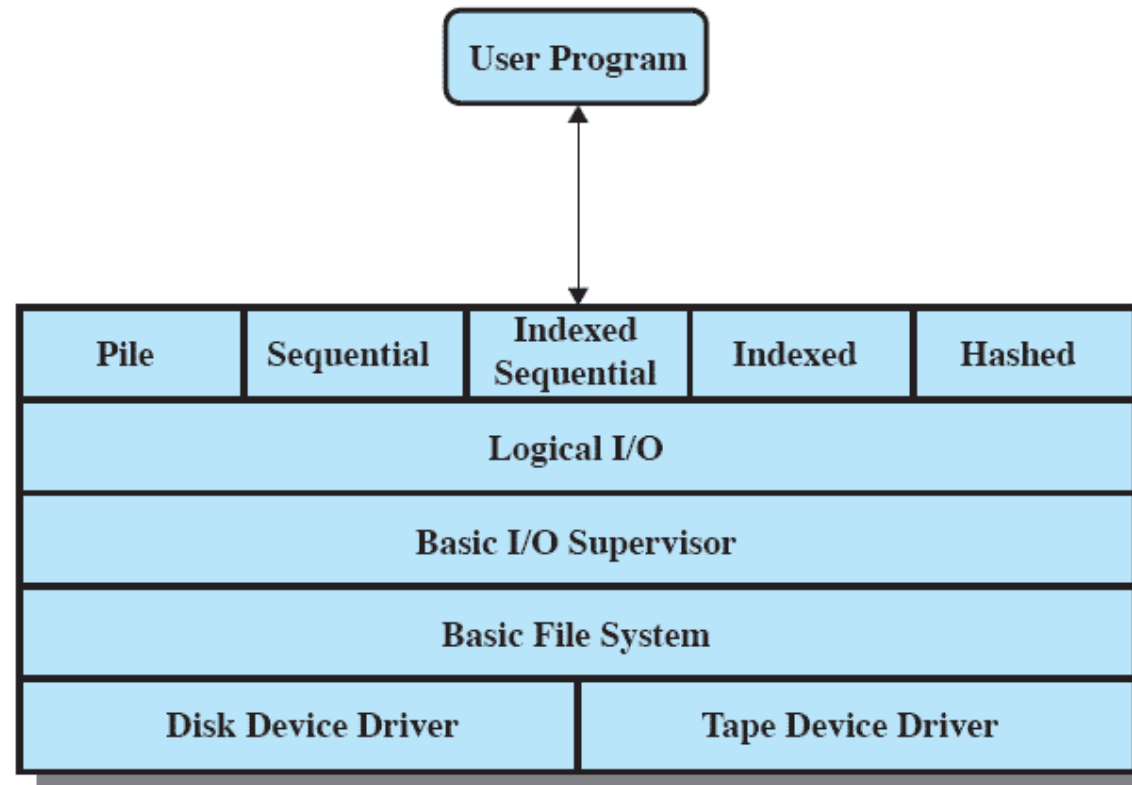


Figure 12.1 File System Software Architecture

Device Drivers

- Lowest level
- Communicates directly with peripherals
- Responsible for starting I/O operations on a device
- Processes the completion of an I/O request
- For file operations, usually either disk or tape drives
- Considered part of the OS

Basic File System

- Physical I/O level, part of OS
- Primary interface with the environment outside the computer
- Deals with exchanging blocks of data
- Concerned with the placement of blocks and buffering blocks in memory
- Does not understand content of the data or structure of the files

Basic I/O Supervisor

- Responsible for all file I/O initiation and termination
- Control structures to deal with:
 - Device I/O
 - Scheduling
 - File status
- Selects and schedules I/O with the device

Logical I/O

- Enables users and applications to access records
- Provides general-purpose record I/O capability
- Maintains basic data about the file

Access method

- Closest to the user
- Reflects different possible file structures
- Provides a standard interface between applications and the file systems and devices that hold the data
- Access method varies depending on the ways to access and process data for the device

File management

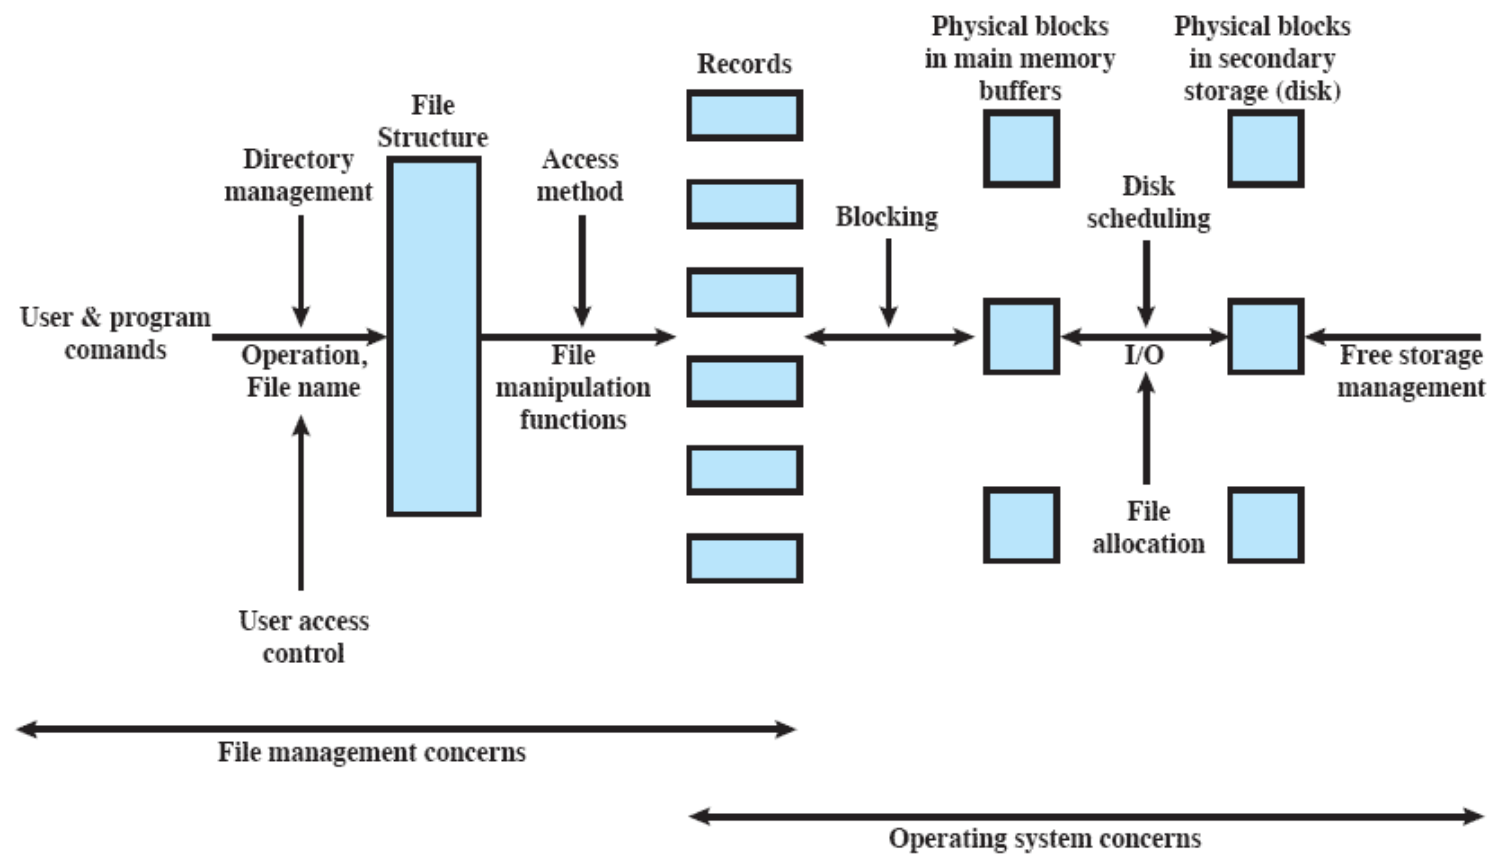


Figure 12.2 Elements of File Management

File Organization

- File organization refers to the logical structure of records
 - Not how they are physically stored
 - Will talk about physical organization later
- Determined by the way in which files are accessed
- So what would be desirable properties of a file organization?

File Organization


- Important criteria include
 - Short access time
 - Ease of update
 - Economy of storage
 - Simple maintenance
 - Reliability
- Priority of these concerns depends on the use (read-only CD vs hard drive)
 - They could even directly conflict

File Organization

- Many different file organizations, but usually variations of the following:
 - Pile
 - Sequential file
 - Indexed sequential file
 - Indexed file
 - Direct (or hashed) file

The Pile

- Data is collected in the order it arrives
 - No structure
- Purpose is to accumulate a mass of data and save it
- Records may have different fields
- Record access is by exhaustive search



Variable-length records
Variable set of fields
Chronological order

(a) Pile File

Sequential File

- Fixed format used for records
- Records are the same length
- All fields are the same (order and length)
- Field names and lengths are attributes of the file
- Key field
 - Uniquely identifies the record
 - Records are stored in key sequence

Fixed-length records
Fixed set of fields in fixed order
Sequential order based on key field

(b) Sequential File

Sequential File


- Performs well on batch operations
- Poorly on individual updates or additions
- Can have separate log or transaction file
 - Periodically merge this file with master
- Alternatively organize file as linked list

Fixed-length records
Fixed set of fields in fixed order
Sequential order based on key field

(b) Sequential File

Indexed Sequential File


- Maintains key characteristics of sequential file
 - Records are organized in sequence by key
- Adds some additional features
 - Index to the file to support random access
 - Overflow file



(c) Indexed Sequential File

Indexed File

- Uses multiple indexes for different key fields
 - May contain exhaustive index
 - One entry for every record in main file
- When a new record is added to main file
 - All of the index files must be updated



(d) Indexed File

Direct or Hashed File

- Lets us access directly any block of a known address
- Key field required for each record
- Hash tells us where it is stored in the overall file
- Need fixed length records

B-Trees

- If file is large single sequential file of indexes does not provide rapid access
 - Need a more structured index file
- Do it in two levels
 - Original file broken into sections
 - Upper level consists of sequenced pointers to the lower-level sections
 - Can extend this to more levels for a tree structure

B-Trees


- Searching a tree structure is only efficient if we have structure
 - If some are small and others are large, time to search index is uneven
- Ensure a balanced tree structure
 - All branches of equal length
 - Self-balancing tree structure
- B-tree is standard way to organize indexes for databases and OS

B-Trees

- B-trees have the following characteristics
 - Tree consists of nodes and leaves
 - Each node contains at least one key that uniquely identifies a record
 - Can contain more than one pointer to child nodes or leaves
 - Each node is limited to a maximum number of keys
 - Keys in a node are nondecreasing
- Advantage to B-trees is very shallow, so relatively low depth

B-Trees


- Might see the *order* of a B-tree
 - Usually defined as the minimum number of keys in a non-root node
 - Knuth defines it as maximum number of children ($\text{max keys} + 1$)




A B-tree of order 5 (Knuth)

Inserting to B-tree


- As if inserting to binary tree, search for key
 - If room in leave node, insert it
- If no room, split it up around median key in node
 - Move median node into parent and hang the now two leaves around it
- If no room for median node in parent, split parent around its median
 - Continue the process to the root if necessary




(a) B-tree of minimum degree $d = 3$.



(b) Key = 90 inserted. This is a simple insertion into a node.



(c) Key = 45 inserted. This requires splitting a node into two parts and promoting one key to the root node.



(d) Key = 84 inserted. This requires splitting a node into two parts and promoting one key to the root node. This then requires the root node to be split and a new root created.

Figure 12.5 Inserting Nodes into a B-tree

Inserting to B-tree

- What if we don't want to have to work our way back up?
- Can perform operations on B-tree always going down
 - Just cannot wait until we find a full parent
 - Instead, while searching to find a new key insertion point
 - Split each full node along the way

File Structure

- At user level we have records, etc
- At lowest level, have alternatives:
 - Byte sequence or stream
 - Maximum flexibility
 - Sequence of records
 - More closely matching user level
- Tree
 - Faster searches

Blocks and records

- Records are the logical unit of access of a structured file
 - But blocks are the unit for I/O with secondary storage
 - Mismatch
- Three approaches are common:
 - Fixed length blocking
 - Variable length spanned blocking
 - Variable length unspanned blocking



Blocks and records

- Tradeoff
 - Larger block size means more records in one read
 - Good if records are next to each other
 - Larger blocks can result in extra records that we don't need
 - Assuming not strong locality of reference
 - Also requires larger buffers

Fixed Blocking

- Fixed-length records are used
 - Integral number of records are stored in a block
- Tends to be default for sequential files with fixed length records
- Unused space at the end of a block is internal fragmentation


Fixed Blocking



Variable Length Spanned Blocking

- Can use variable-length records
 - Could then pack into blocks with no unused space
- Would need to allow records to span multiple blocks
 - At end of block, need a pointer to the next one
- More efficient in terms of storage
- Difficult to implement
 - Records that span two blocks require two I/O operations
 - Difficult to update records

Variable Length Spanned Blocking



Variable Blocking: Spanned



Data



Gaps due to hardware design



Waste due to block fit to track size



Waste due to record fit to block size




Waste due to block size constraint from fixed record size


Variable-length Unspanned Blocking

- Variable length records
 - This time do not allow them to span multiple blocks
- Wasted space in most blocks, as cannot use remainder
- Also cannot have record larger than a block

Variable-length Unspanned Blocking



Variable Blocking: Unspanned



Filenames

- Files usually accessed by name
 - File types should be known by the OS
 - Should not print binary files, try and compile an image, etc
 - Historically done with file extensions
 - `file` command in UNIX examines magic number in file
- Files can also be accessed by a link

File links

- Links can be either `hard` or `soft` links
- Hard links are additional alias for file
 - Two or more filenames for the same physical thing
 - If you delete one hard link (`rm`), not deleted until all hard links `rm`'ed
 - Share the exact same physical blocks
 - Can only exist on the same filesystem
 - Cannot create hard link to a directory
 - Could get stuck in endless cycle

File links

- Soft links are just a pointer to a filename
 - Treated like the file it is pointing at
 - Distinguishable from the file
 - May point to non-existent files
 - Files on other systems
 - Can only exist on the same filesystem

File types in Unix

- Regular files
 - Most common type
 - Treated as a byte stream, no kernel support of structure
- Directories
 - Binary file containing list of files in it
 - Each entry is file/inode pair
 - Used to associate nodes and directory locations
 - Data itself does not know logical organization

File types in Unix

- Character-special and block-special files
 - Allows applications to communicate with hardware and peripherals
 - In /dev directory
 - Kernel handles links to device drivers
- Character-special files
 - Devices perform their own buffering, raw character stream
- Block-special files
 - Kernel handles buffering

Secondary Storage Management

- Operating system has to allocate blocks to files
 - How do we do this?
- Two related issues
 - Space must be allocated to files
 - Must keep track of the space available for allocation

File Allocation Issues

- When a file is created, is maximum space allocated at once?
- If space is added to a file in contiguous chunks, what should be the size of the portion?
 - What data structure should be used to keep track of the file portions?
 - FAT or inode

File Allocation Issues

- Preallocation
 - Needs maximum size for file at time of creation
 - We don't have a reliable oracle, so tough to do in practice
 - Tends to result in overestimation at best

Portion Size

- Two extremes:
 - Portion large enough to hold entire file is allocated
 - Allocate space one block at a time
- Trade-off between efficiency from the point of view of:
 - A single file
 - Overall system efficiency

File Allocation Method

- Three methods in common use:
 - Contiguous
 - Chained
 - Indexed

Contiguous Allocation

- Single set of blocks is allocated to a file at the time of creation
- Only a single entry in the file allocation table
 - Starting block and length of the file
- External fragmentation will occur
 - Need to perform compaction

Contiguous File Allocation

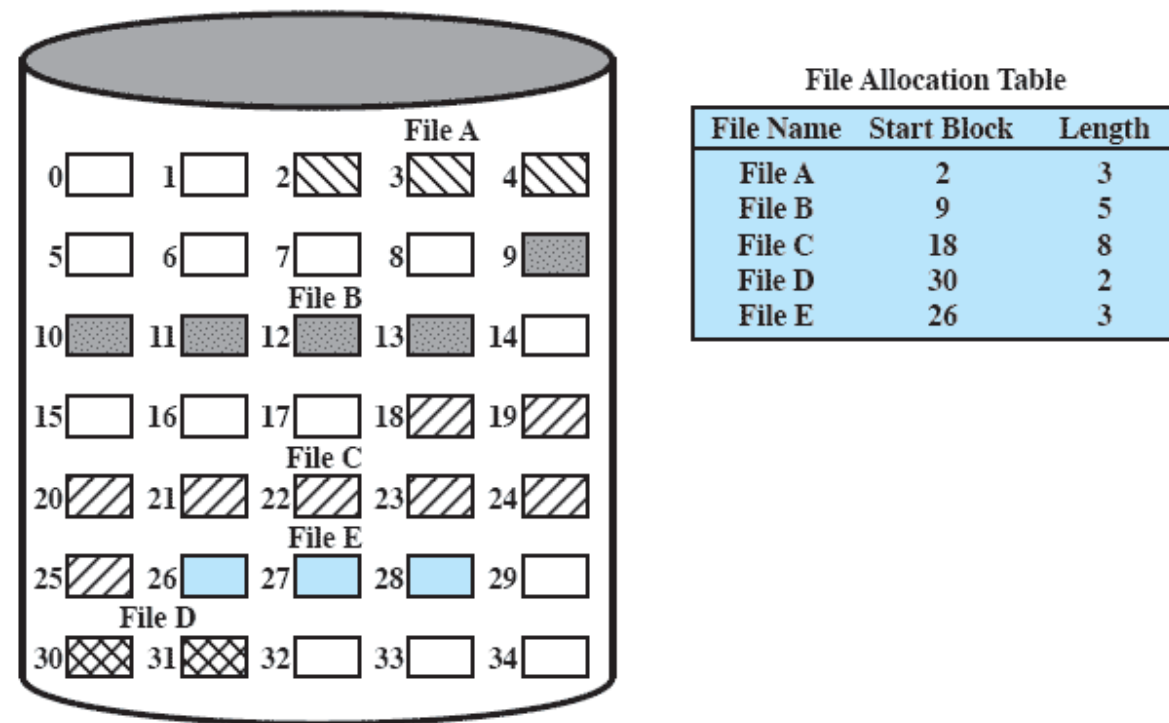


Figure 12.7 Contiguous File Allocation

External Fragmentation

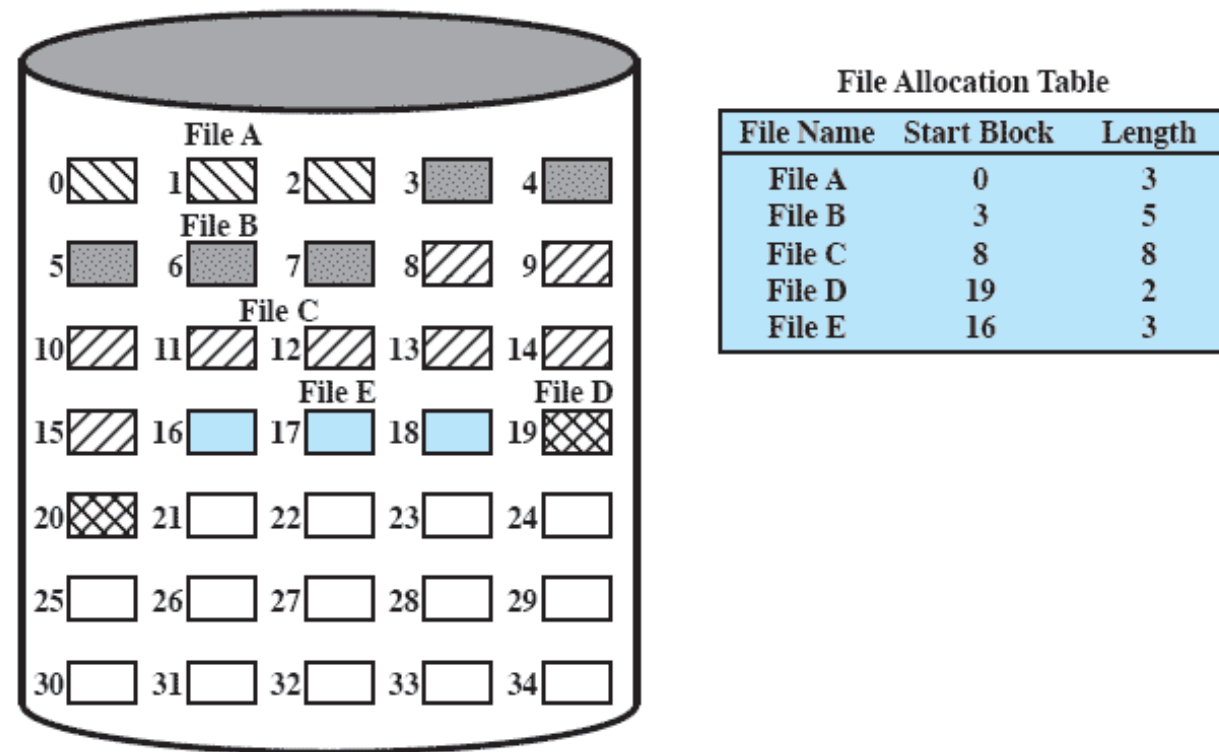


Figure 12.8 Contiguous File Allocation (After Compaction)

Chained Allocation

- Allocation on basis of individual block
- Each block contains a pointer to the next block in the chain
- Only single entry in file allocation table
 - Starting block and length of file
- No external fragmentation
- Best for sequential files

Chained Allocation




Figure 12.9 Chained Allocation

Chained Allocation

- How efficient is this from a physical perspective?
- Reading requires constant seeks
- Can consolidate to help with this

Chained Allocation Consolidation




Figure 12.10 Chained Allocation (After Consolidation)

Indexed Allocation

- File allocation table contains separate one-level index for each file
- The index has one entry for each portion allocated to the file
- The file allocation table contains block number for the index

Indexed Allocation Method

- Allocation may be either:
 - Fixed size blocks
 - Variable sized blocks
- Allocating by blocks eliminates external fragmentation
- Variable sized blocks improves locality
- Both cases require occasional consolidation
- Important consideration is that this index needs to be saved in a block!
 - Takes up space, which could be a problem if we want really large files
 - What happens if block size isn't large enough to hold index?

Indexed Allocation with Block Portions

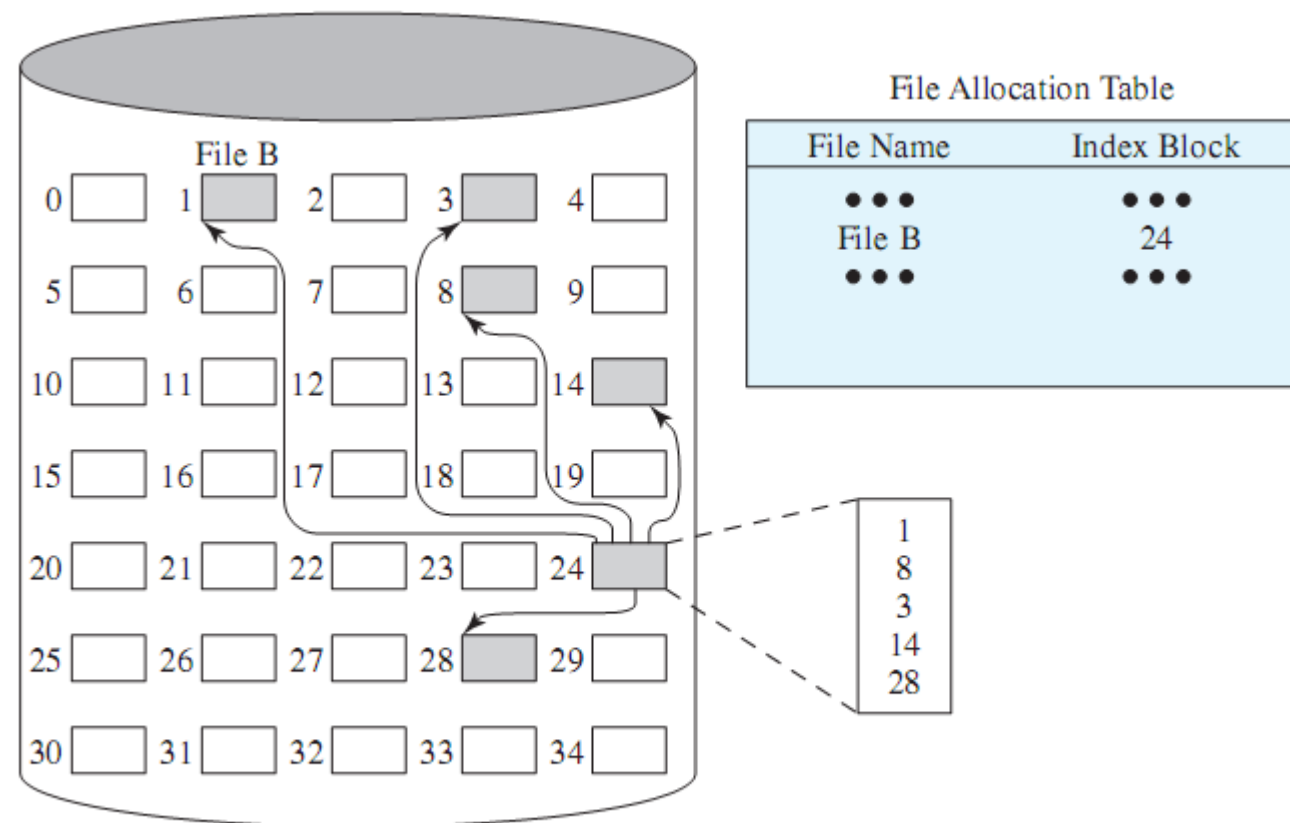


Figure 12.11 Indexed Allocation with Block Portions

Indexed Allocation with Variable Length Portions

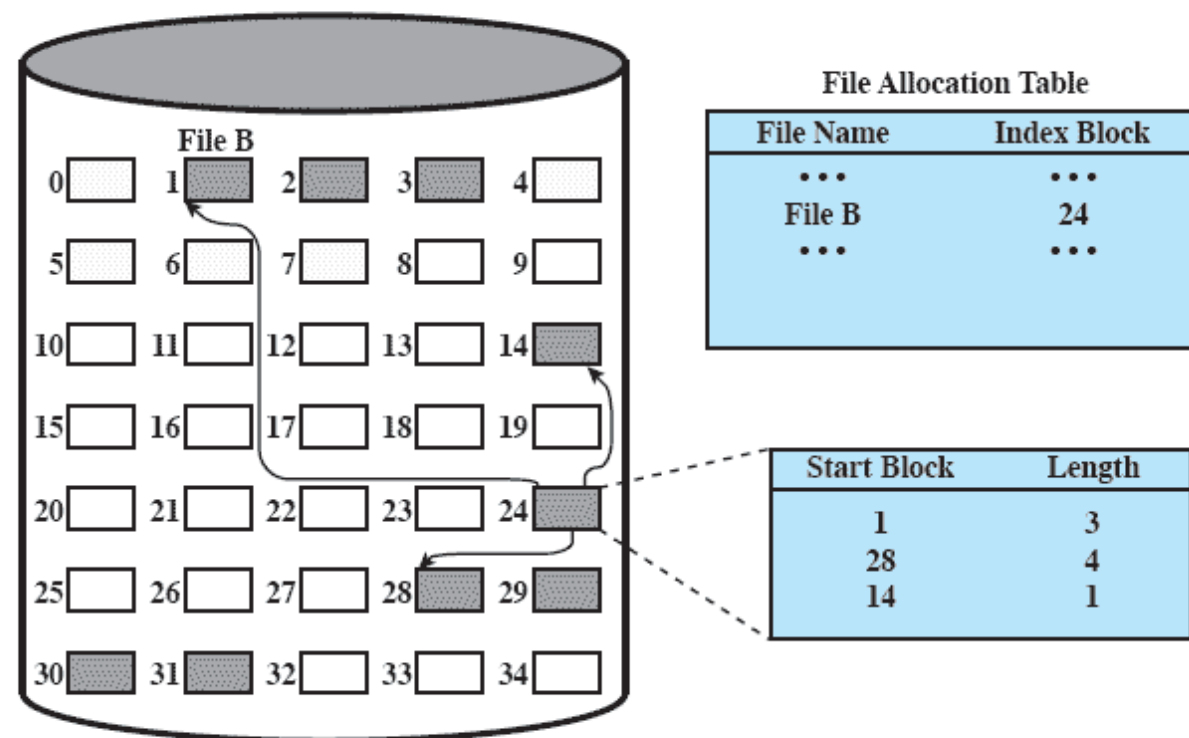


Figure 12.12 Indexed Allocation with Variable-Length Portions

Inodes

- Index node
- Control structure that contains key information for a particular file
- Several filenames may be associated with a single inode
 - But an active inode is associated with only one file
 - Each file is controlled by only one inode

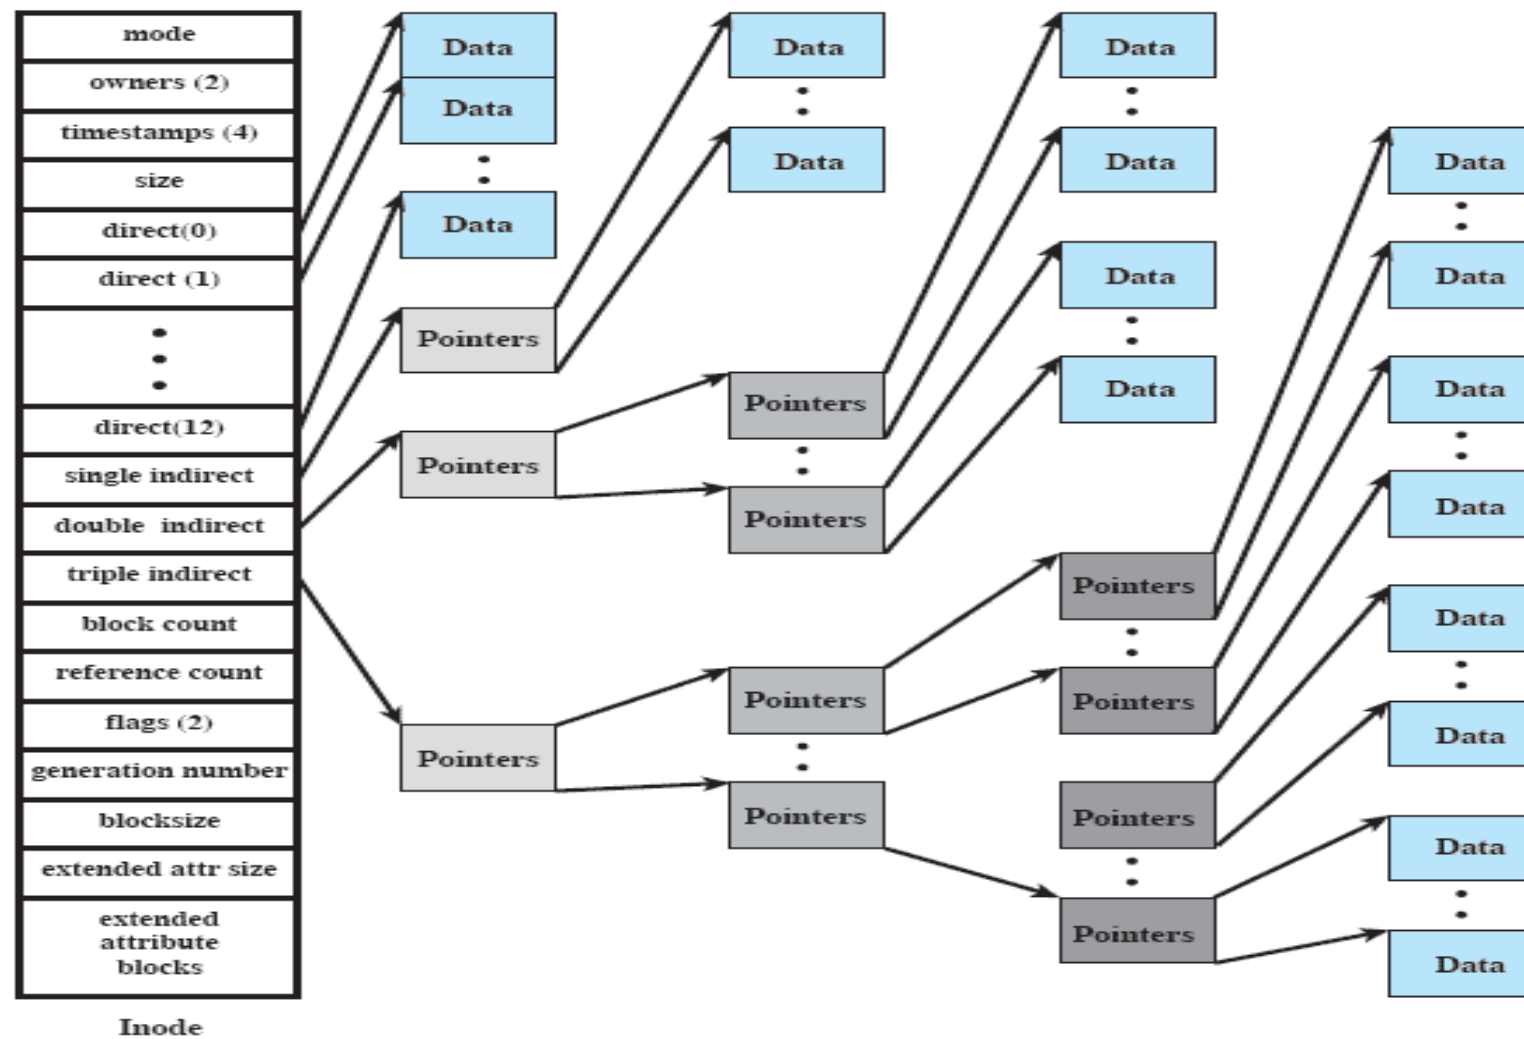
FreeBSD Inodes include:

- The type and access mode of the file
- The file's owner and group-access identifiers
- Creation time, last read/write time
- File size
- Sequence of block pointers
- Number of blocks and Number of directory entries
- Blocksize of the data blocks
- Kernel and user settable flags
- Generation number for the file
- Size of Extended attribute information
- Zero or more extended attribute entries

Inodes

- File allocation is done on a block basis
- Allocation is dynamic
 - Blocks may not be contiguous
- Index method keeps track of files
 - Part of index stored in the file inode
- Inode includes a number of direct pointers
 - Three levels of indirection

Inodes



Inodes

- File allocation is done on a block basis
- Allocation is dynamic
 - Blocks may not be contiguous
- Index method keeps track of files
 - Part of index stored in the file inode
- Inode includes a number of direct pointers
 - Three levels of indirection

Inodes

- This level of indirection determines the maximum blocks a file can contain
 - Yet still allows us to access particular blocks efficiently
- Multiple levels lets us access small files without indirection
- Inode is fixed in size so can be kept in memory for long periods of time
- Theoretical maximum file size is huge

Inodes

- Consider an example:

4Kbytes block size

Each block can hold 512 block addresses

Direct can access 12 blocks	48k
Single indirect, 512 blocks	2M
Double indirect, $512 * 512$ blocks	1G
Triple indirect, $512 * 512 * 512$ blocks	512G

More on file management!

- Any questions?